



Persistent Storage with Kubernetes in Production

Which solution and why?

Cheryl Hung, Product Manager
@oicheryl



Cheryl
Product manager
[@oicheryl](https://twitter.com/oicheryl)



- Why state is tricky in Kubernetes
- How do I evaluate my storage needs?
- How do I choose a storage solution for Kubernetes?



- Why state is tricky in Kubernetes
- How do I evaluate my storage needs?
- How do I choose a storage solution for Kubernetes?

Anti-objective:

- Should I use a database, message queue, NoSQL db... for my app?



Why is state so tricky?



Why do I need storage?





Problem 1: No pet storage

@oicheryl



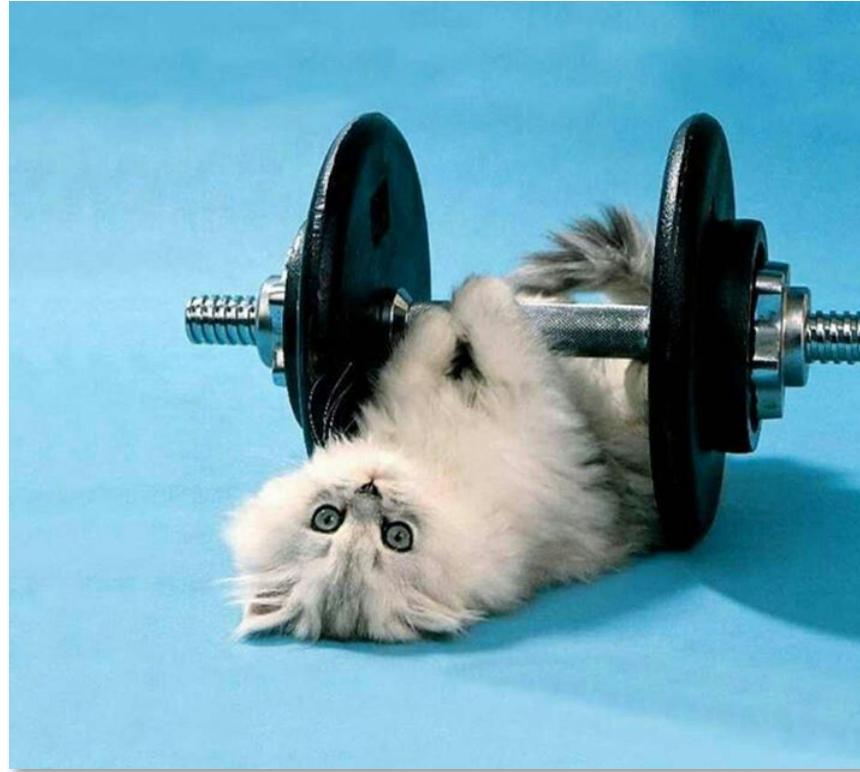
Problem 2: Data needs to follow

@oicheryl



Problem 3: Gravity can be a problem

@oicheryl



Problem 4: Storage isn't just storage

@oicheryl



App Binaries



App data



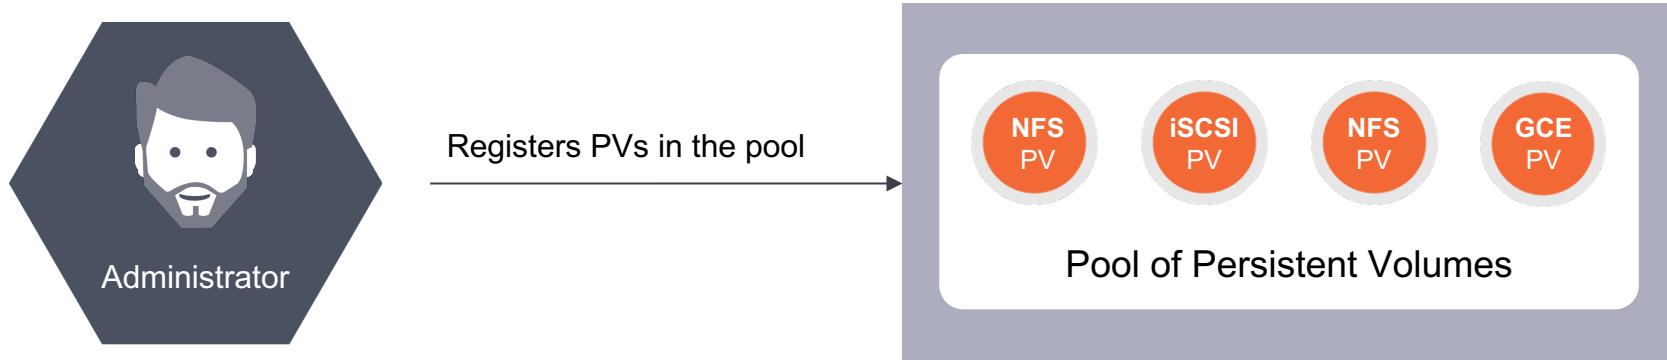
Config



Backup

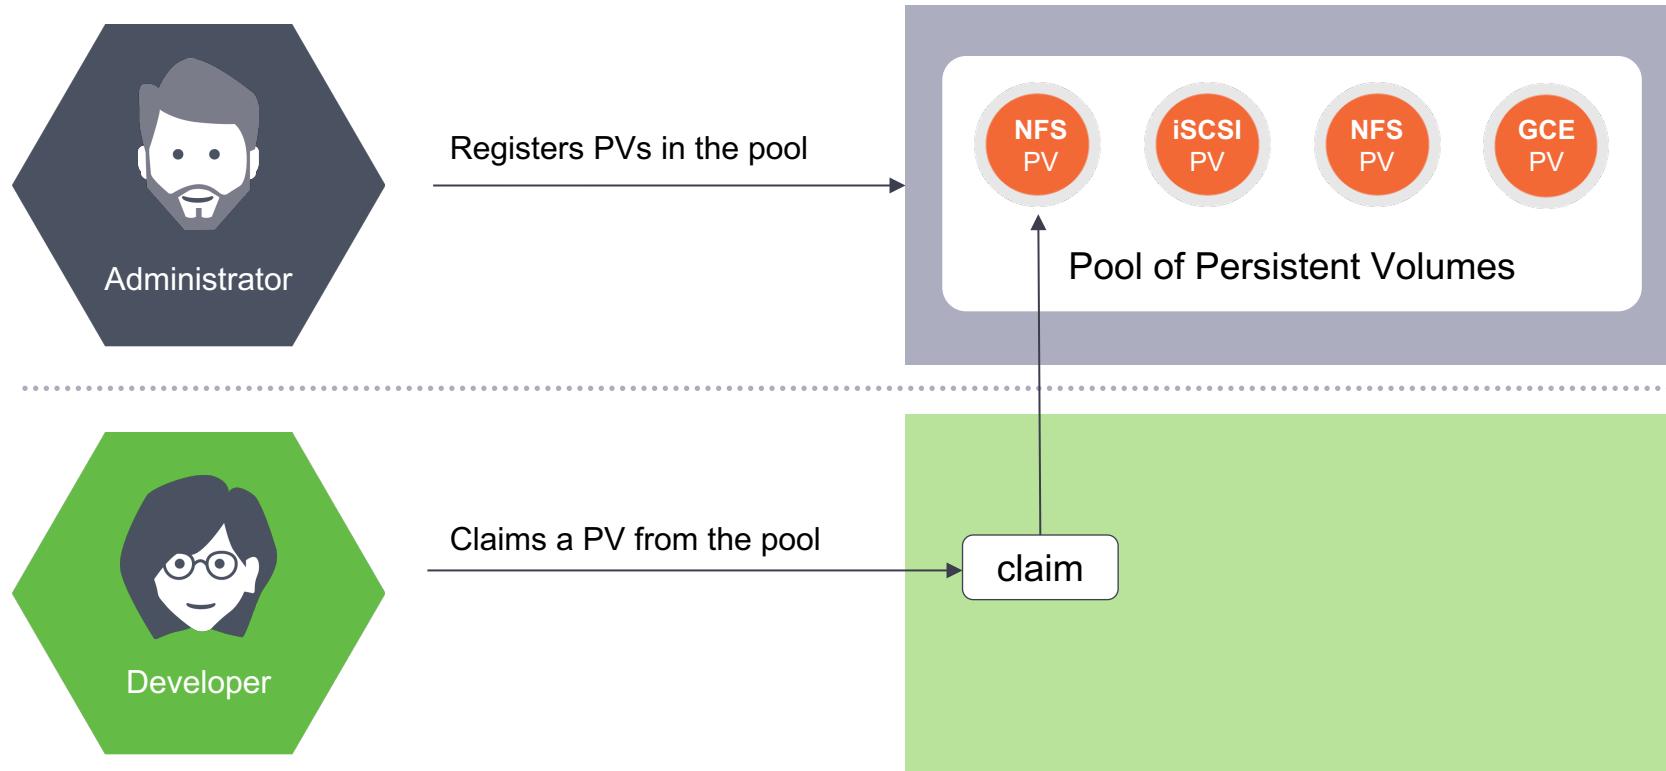
Kubernetes Storage Model: Persistent Volumes and Claims

@oicheryl



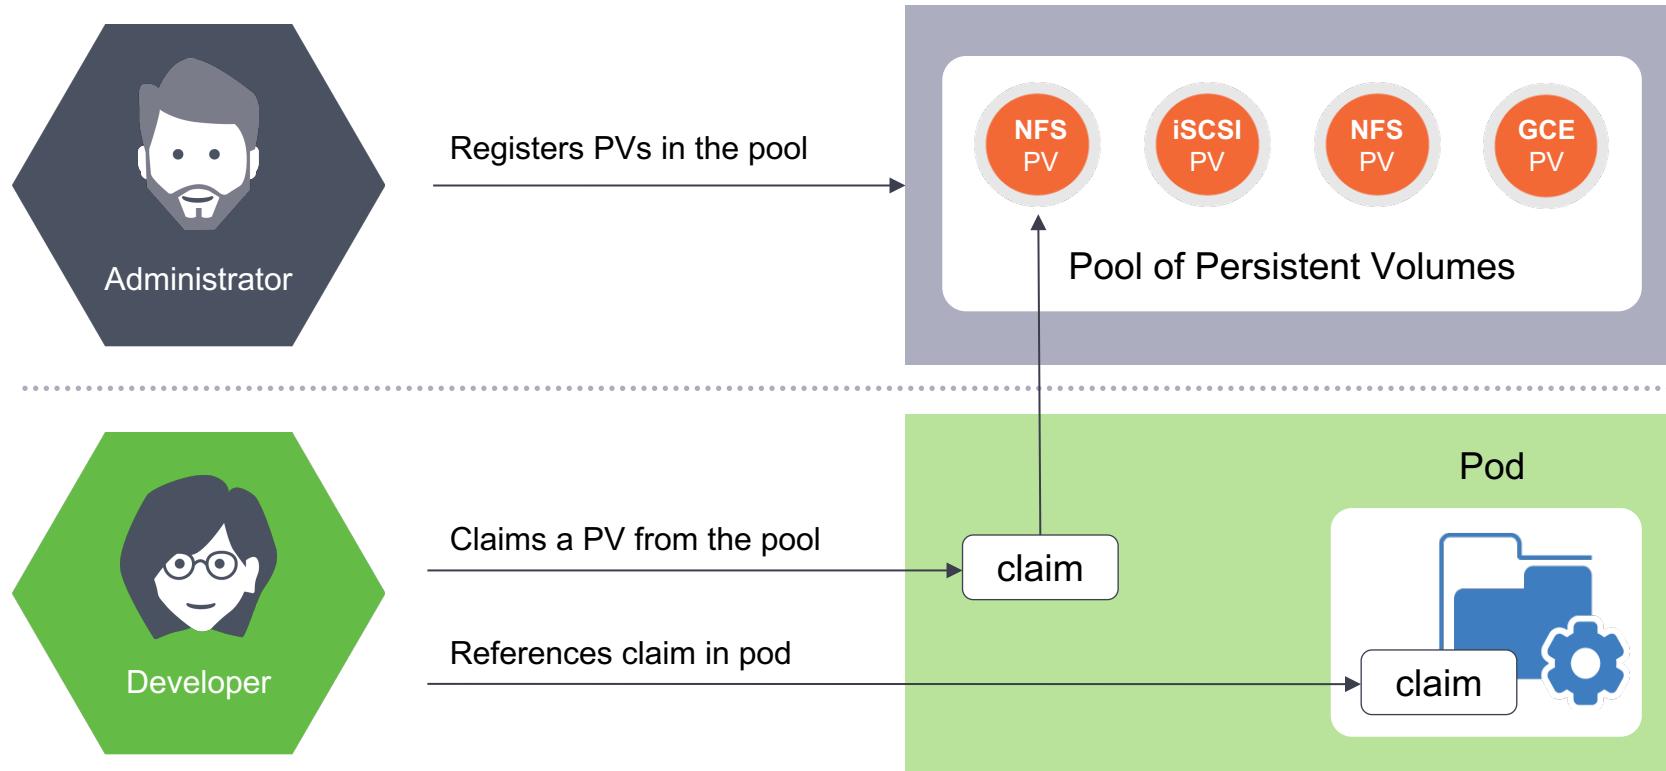
Kubernetes Storage Model: Persistent Volumes and Claims

@oicheryl



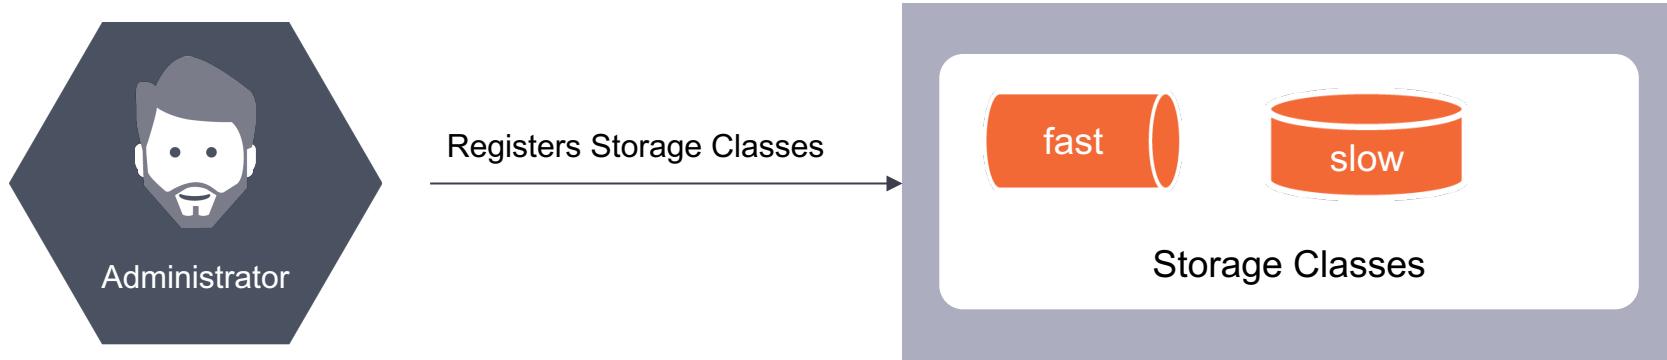
Kubernetes Storage Model: Persistent Volumes and Claims

@oicheryl



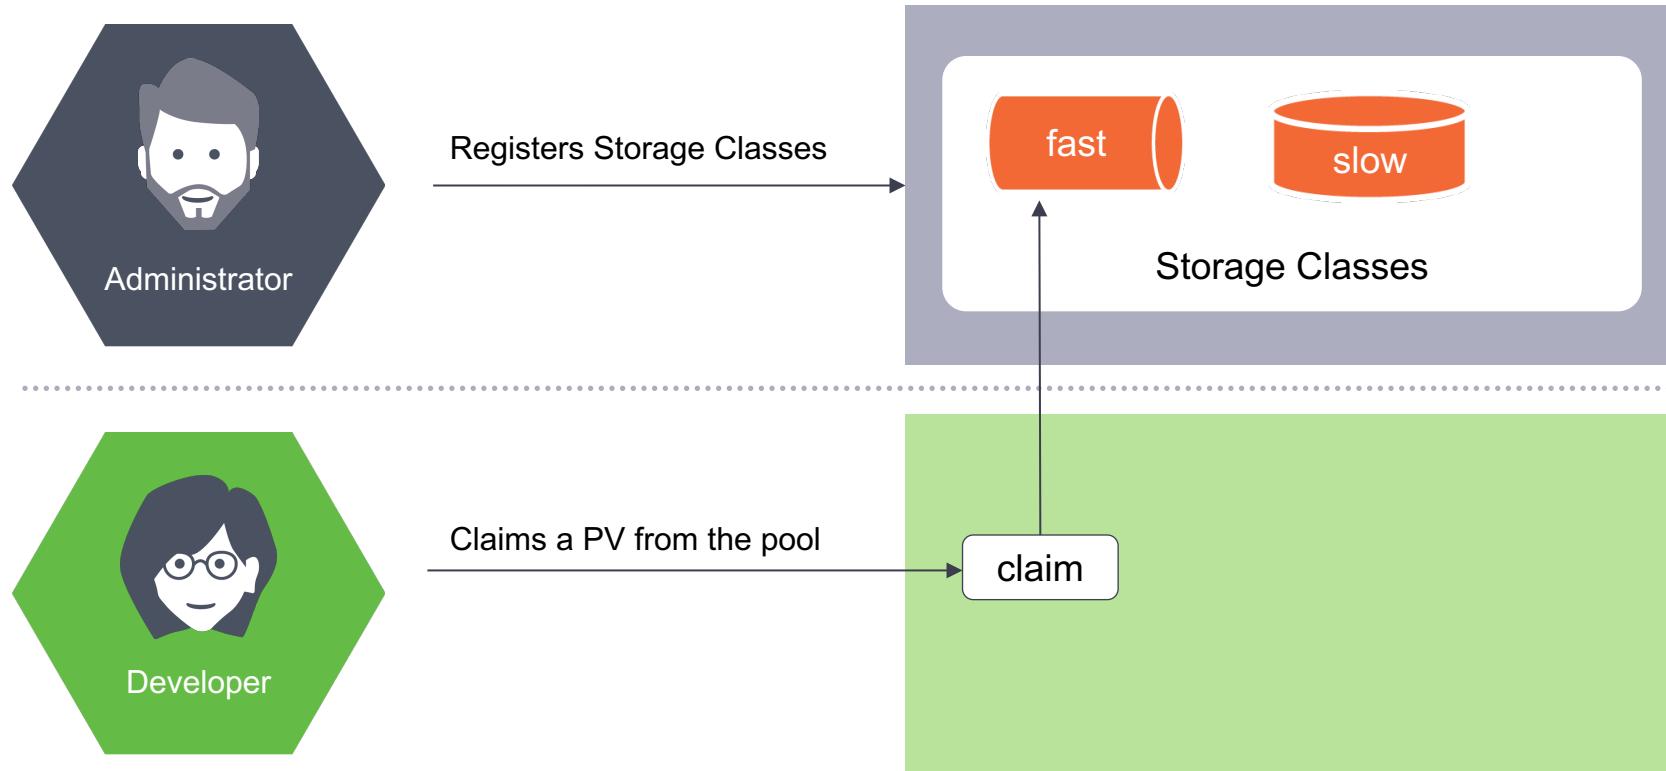
Dynamic provisioning with Storage Classes

@oicheryl



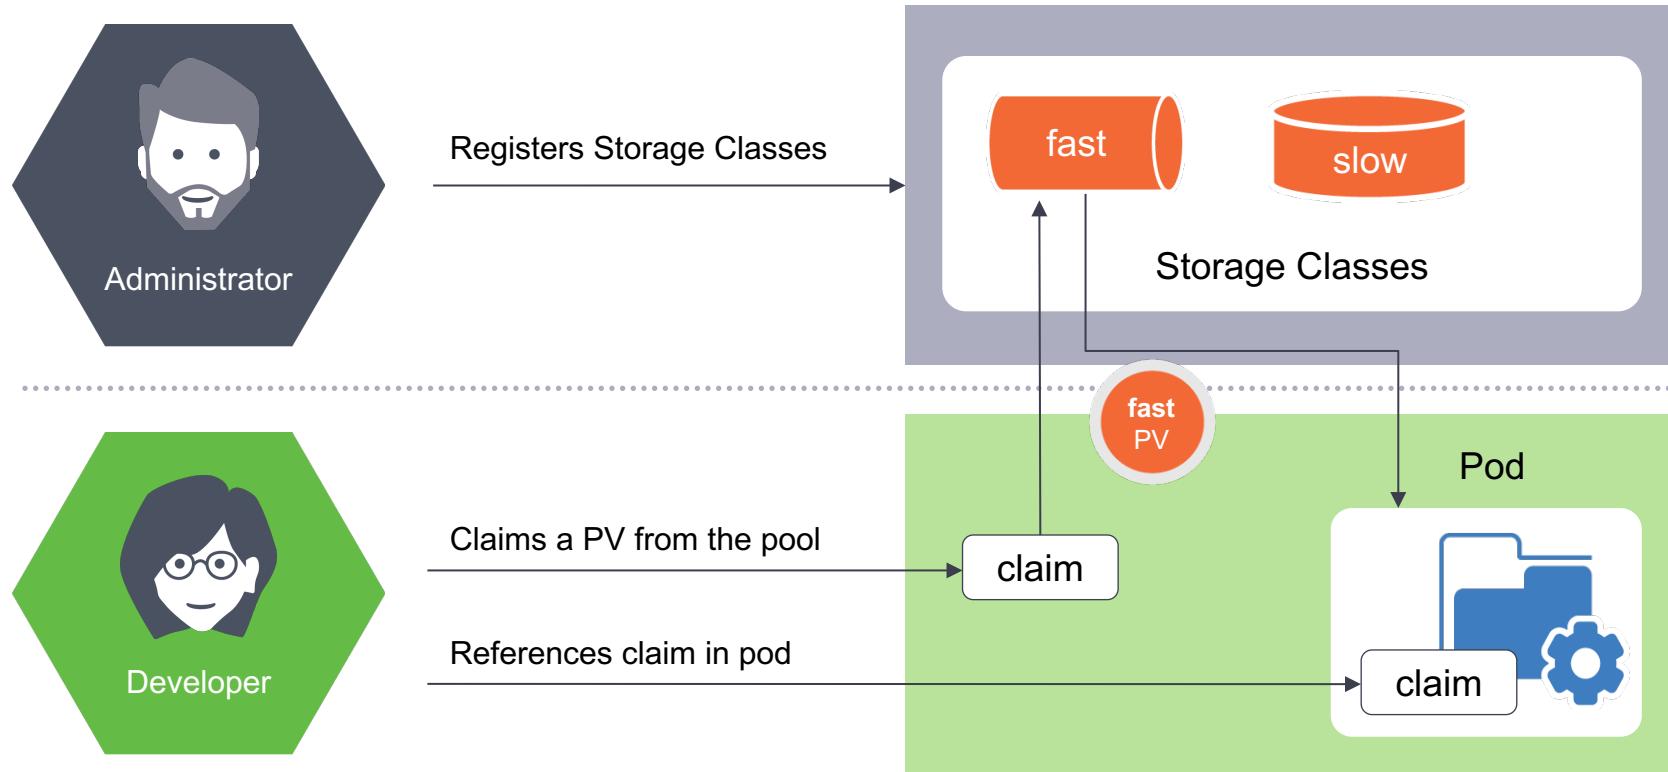
Dynamic provisioning with Storage Classes

@oicheryl



Dynamic provisioning with StorageClass

@oicheryl



Eight Principles of Cloud Native Storage





- A DevOps engineer in a bank
- She wants to migrate a Postgres database to containers

- Horizontally scalable
- No single point of failure
- Resilient and self healing
- Minimal operator overhead
- Decoupled from the underlying platform



1

Storage should be presented to and consumed by applications, not by operating systems or hypervisors.

Storage needs to be able to follow an application as it scales, grows, and moves between platforms and clouds.





The storage platform should be able to run anywhere.
Upgrades and scaling is non-disruptive.



3

Storage resources should be declared and composed just like all other resources required by applications and services.





Storage resources and services should be easy to be provisioned, consumed, moved and managed via an API.



Storage services should integrate and inline security features such as encryption and RBAC.





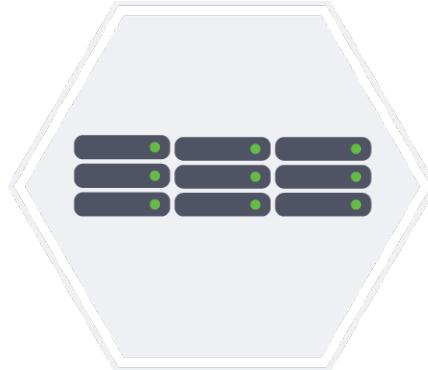
The platform should be able to move application data between locations, dynamically resize volumes for growth, take point in time copies of data for data retention or to facilitate rapid recovery of data.





The storage platform should be able to offer deterministic performance in complex distributed environments.





Block Storage

Data stored in a fixed-size ‘blocks’ in a rigid arrangement – ideal for enterprise databases

File Storage

Data stored as ‘files’ in hierarchically nested ‘folders’ – ideal for active documents

Object Storage

Data stored as ‘objects’ in scalable ‘buckets’ – ideal for unstructured big data, analytics and archiving

8. Consistently available

@oicheryl

8

The storage platform should manage data distribution with a predictable, proven data model to ensure high availability, durability, consistency of application data.

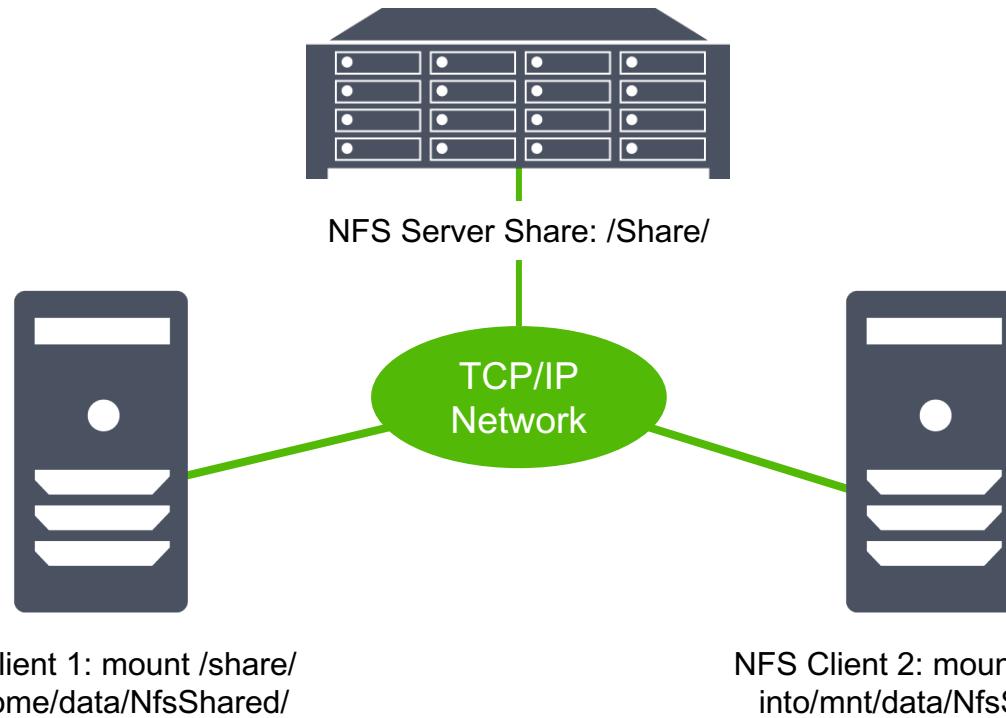


Storage Landscape



Centralized file system

@oicheryl



Single point of failure
Hard to scale horizontally
No native integration

0

Proprietary storage array

@oicheryl

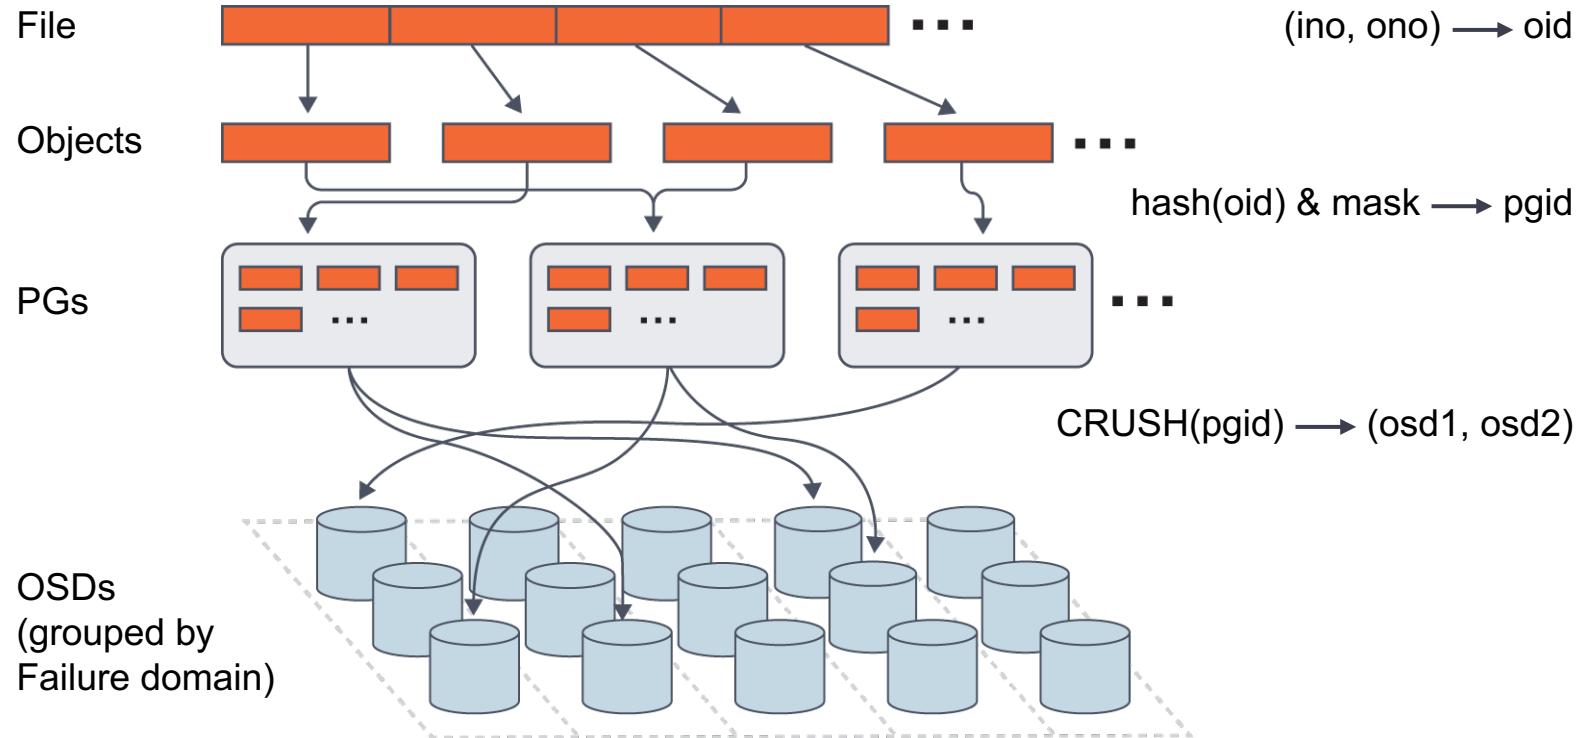


Deterministic performance
Vendor lock in
No thin provisioning
Hard to scale horizontally
Expensive and long lead times

2

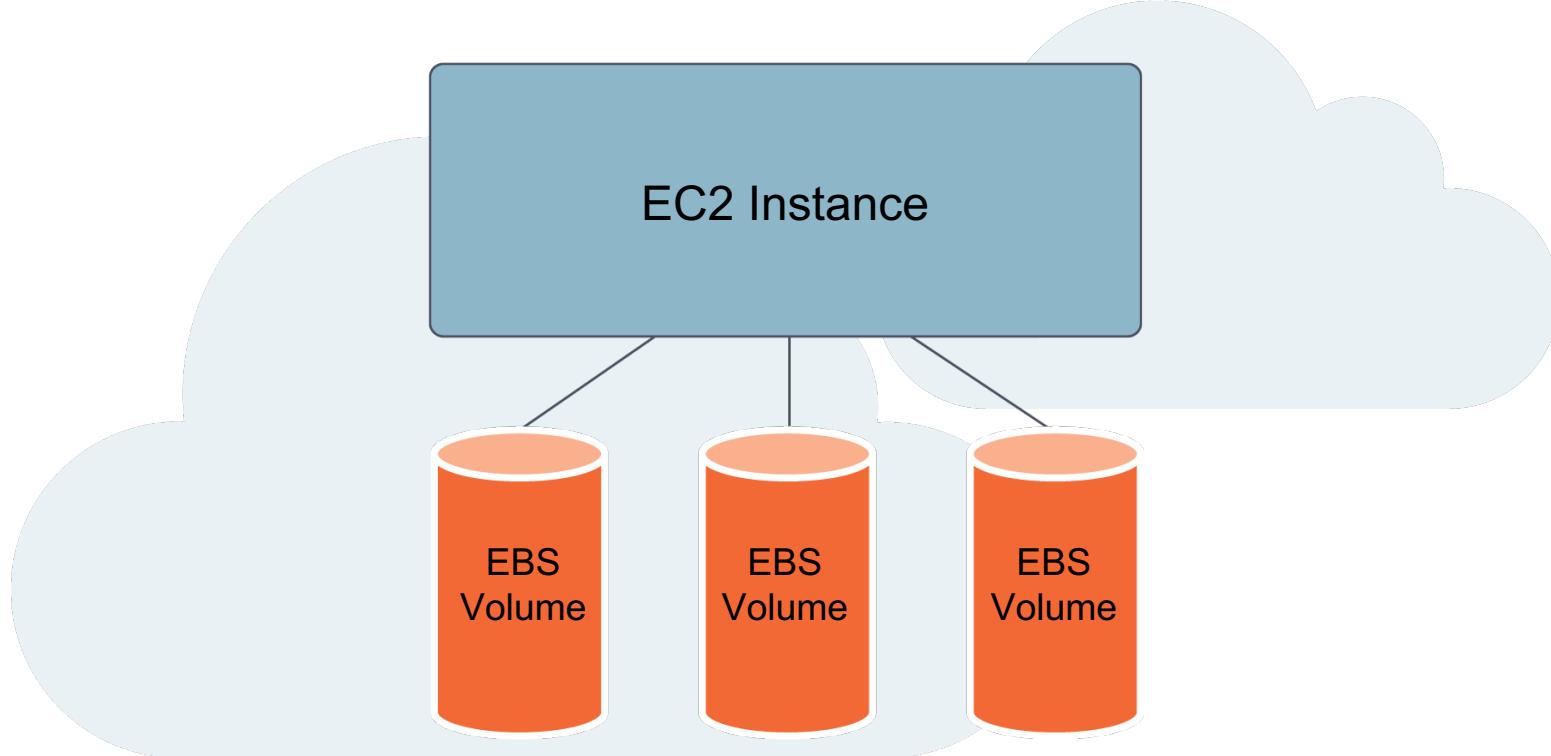
Distributed object store

@oicheryl



- Horizontally scalable
- Hardware agnostic
- Complicated to set up (see: Rook)
- Writes fan out 13-40 times
- Failures are expensive

4



6

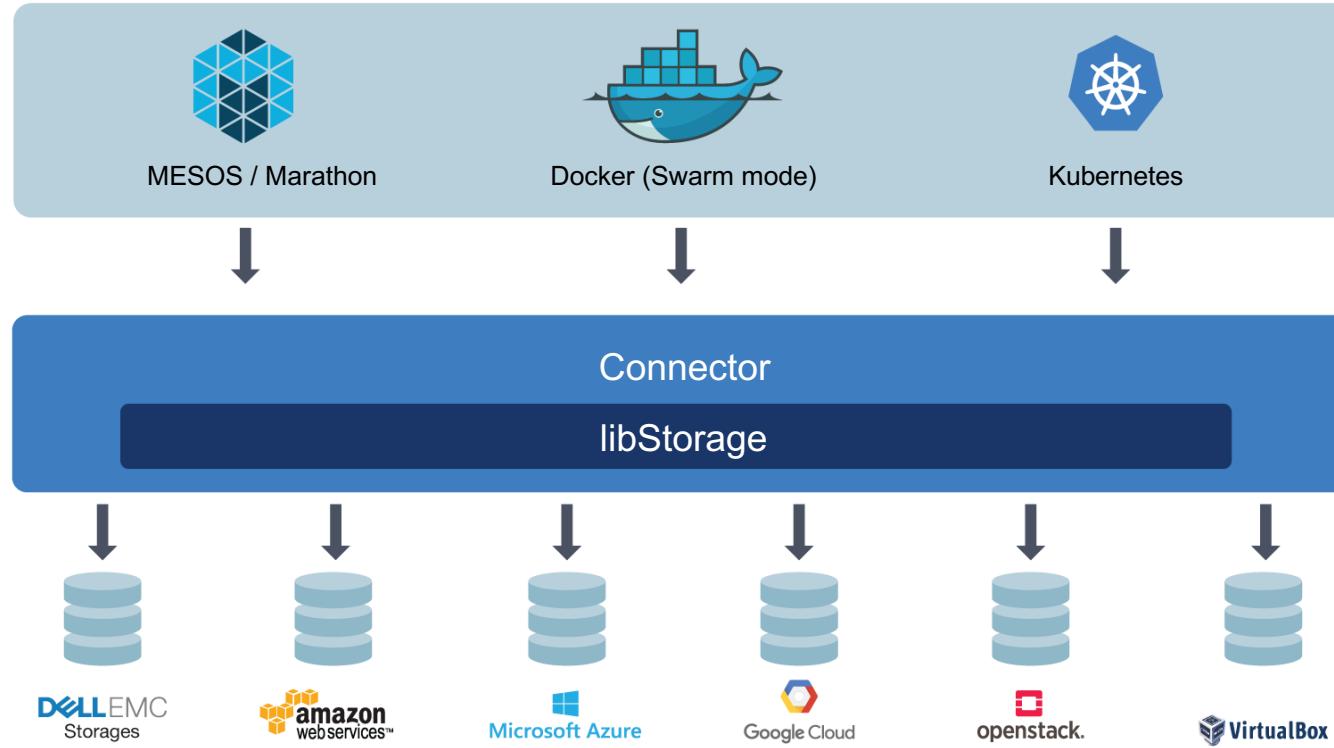
Horizontally scalable

Consistent and performant

40 EBS volumes per EC2 instance

Mounting physical block devices is slow

Expensive, vendor lock in, compliance

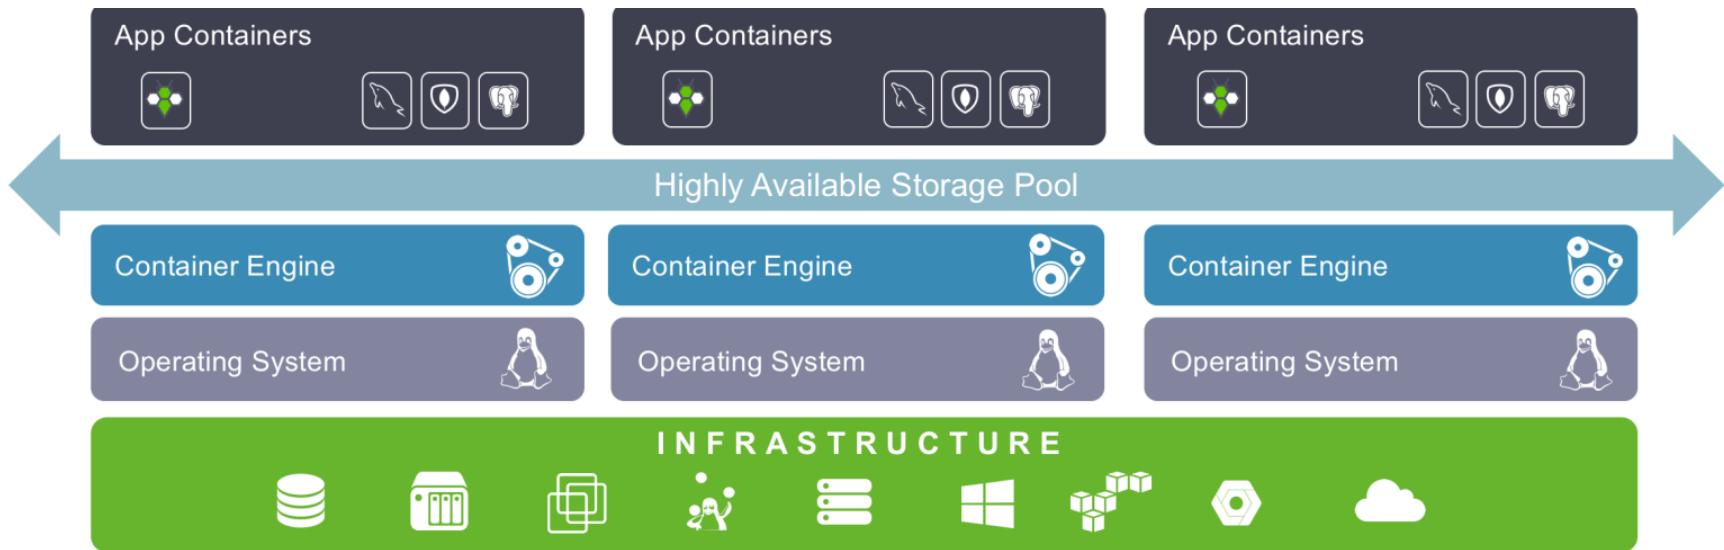


Integrated with Docker/Kubernetes
No guarantees for performance,
replication, failover, scalability

5

Orchestrated storage (StorageOS)

@oicheryl



A software-defined, scale-out storage platform for running enterprise containerized applications in production



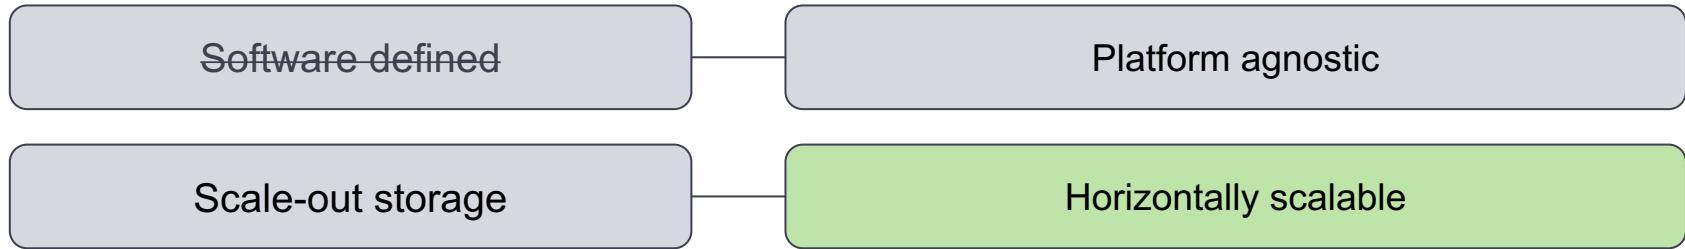
Orchestrated storage (StorageOS)

@oicheryl



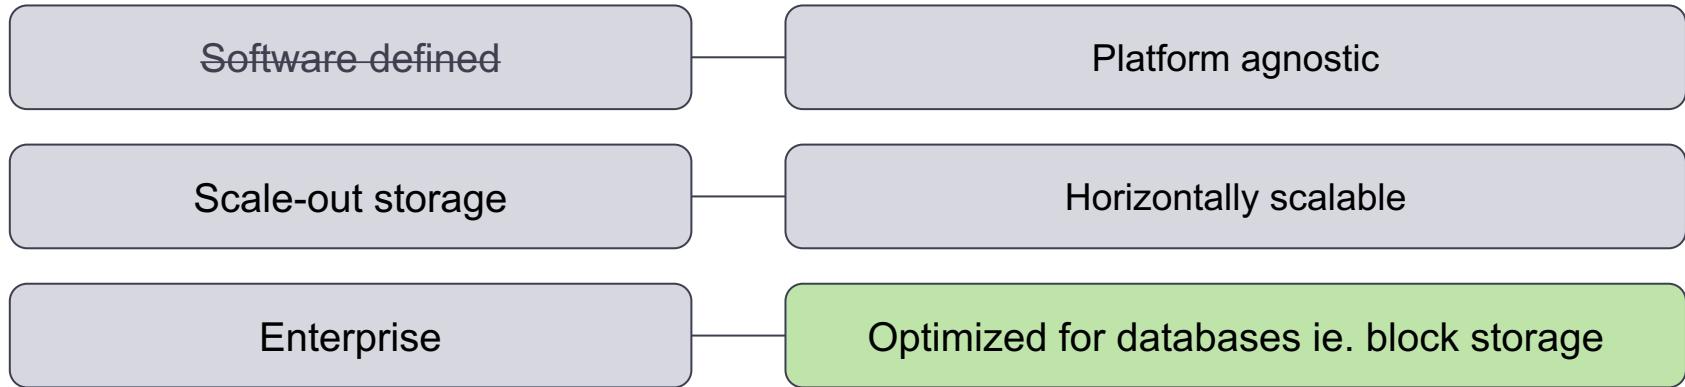
Orchestrated storage (StorageOS)

@oicheryl



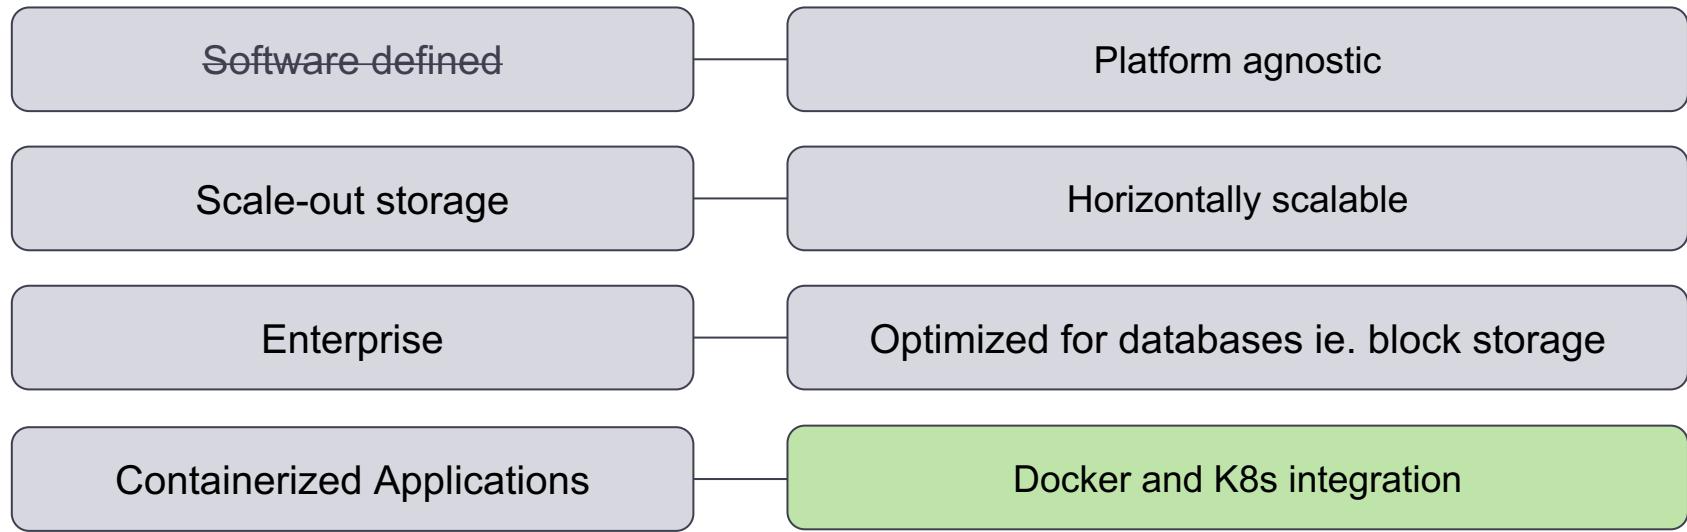
Orchestrated storage (StorageOS)

@oicheryl



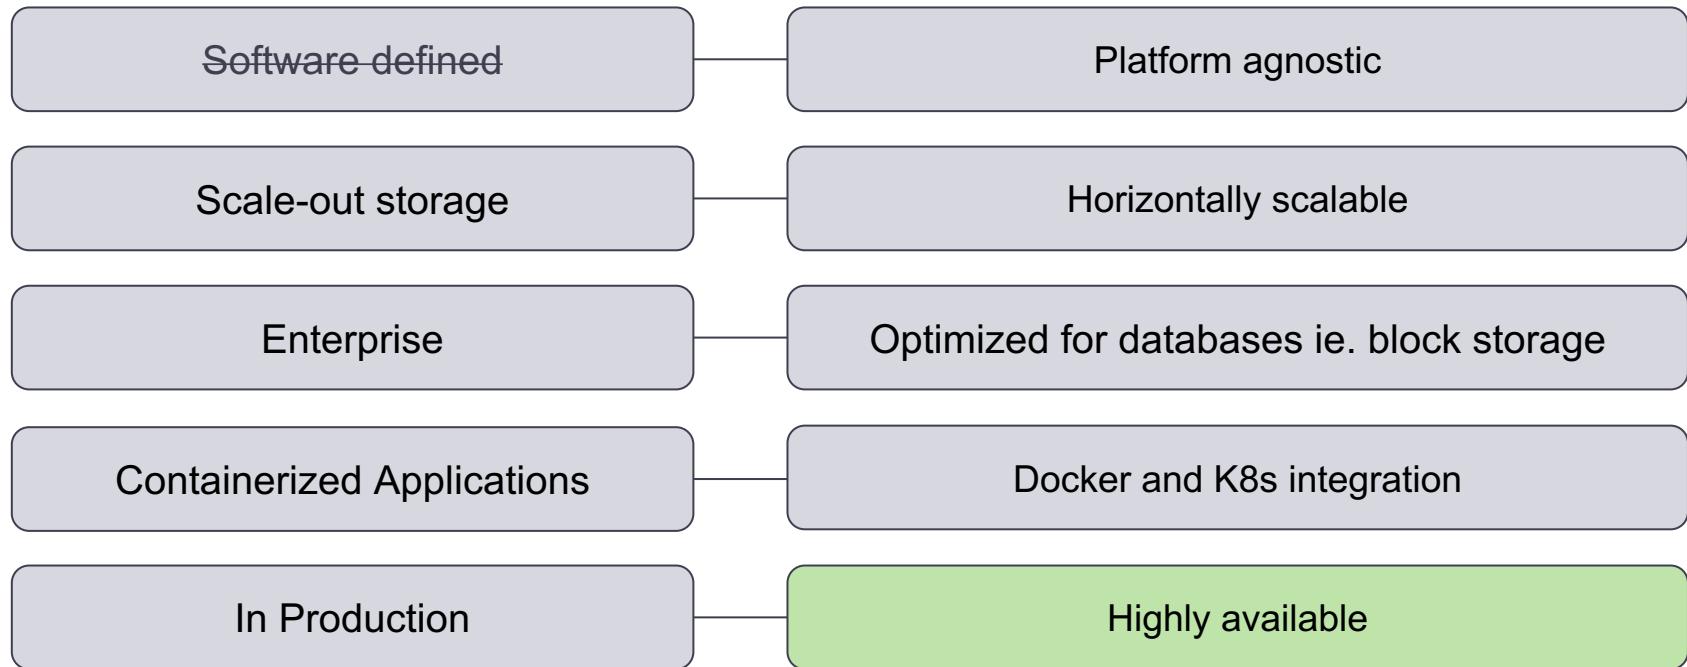
Orchestrated storage (StorageOS)

@oicheryl



Orchestrated storage (StorageOS)

@oicheryl





**KEEP
CALM
IT IS
DEMO
TIME**

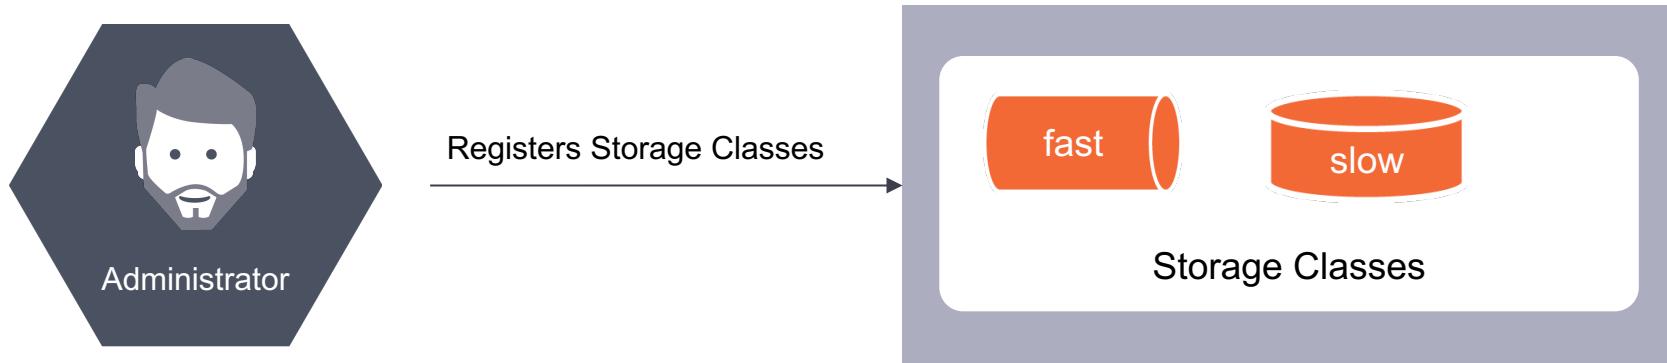
Orchestrated storage (StorageOS)

@oicheryl



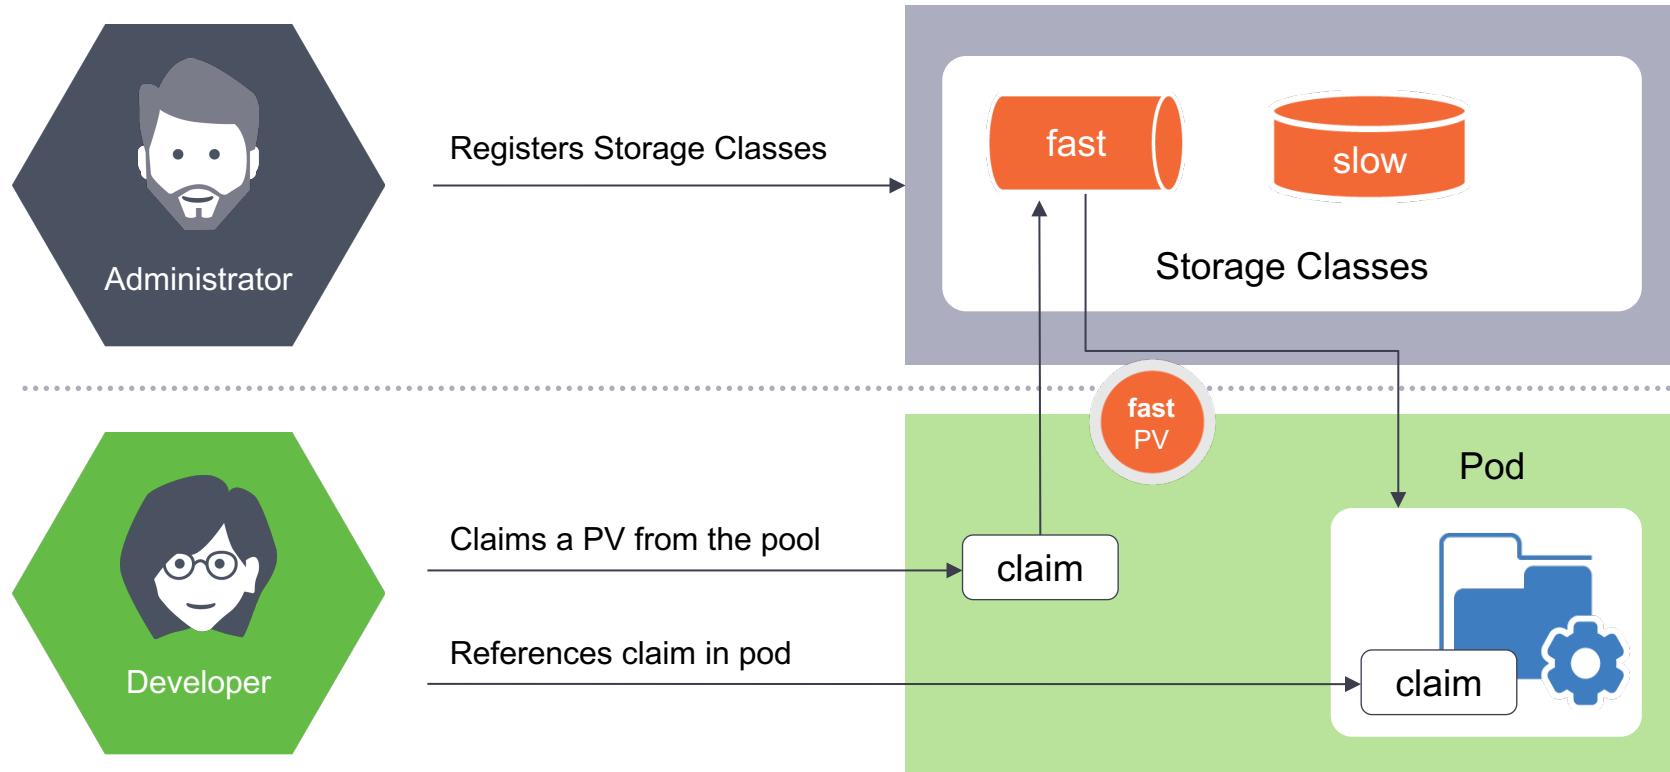
Create a StorageClass

@oicheryl



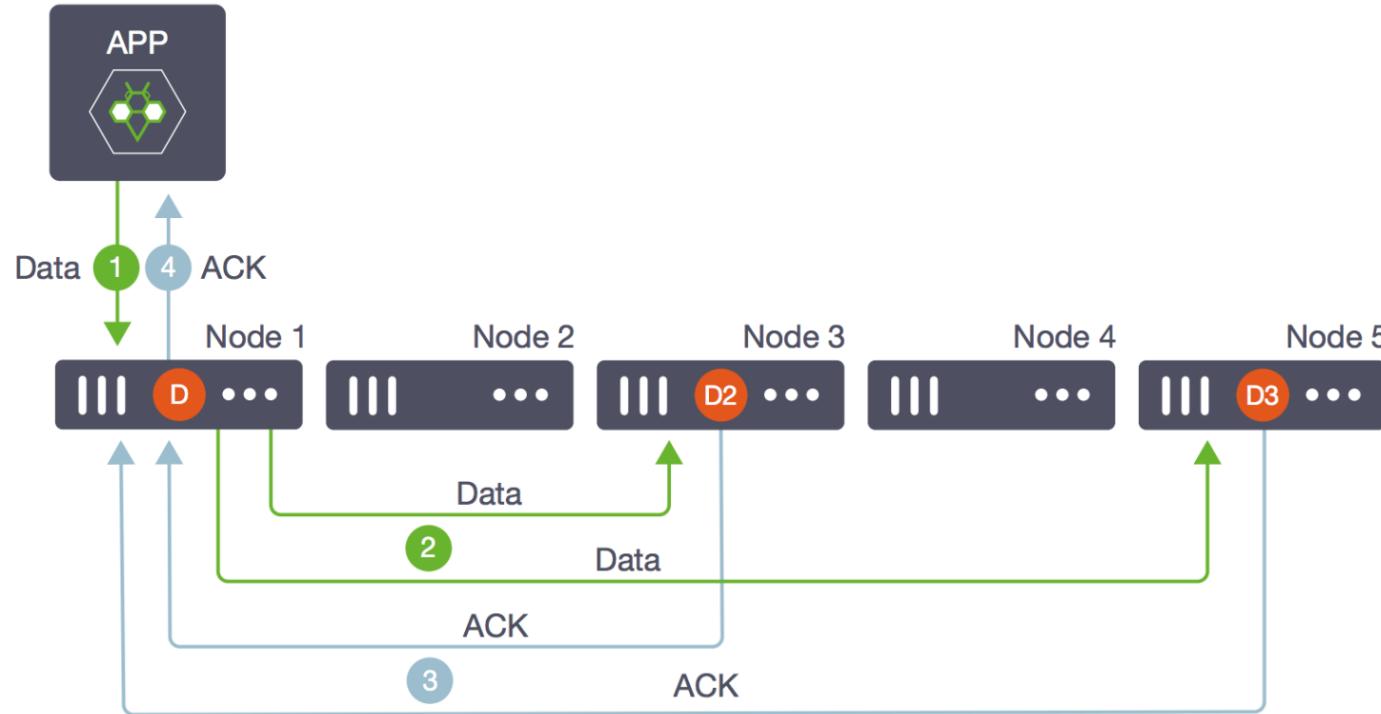
Claim a persistent volume

@oicheryl



High availability with StorageOS

@oicheryl



8

Horizontally scalable
Consistent and performant
Platform agnostic
Synchronous replication
Volume is limited to the size of one node

To Recap...



But you can compare storage using the eight principles:

- | | |
|----------------------------------|---------------------------|
| 1. Application centric | 5. Natively secure |
| 2. Platform agnostic | 6. Agile |
| 3. Declarative and composable | 7. Performant |
| 4. API driven | 8. Consistently available |



Objective is to define an industry standard “Container Storage Interface” (CSI) that will enable storage vendors to develop a plugin once and have it work across a number of container orchestration systems.

StorageOS gives you the **reliability**, **scalability** and **developer ease** of managed service, at the **performance** and **price** of local storage.





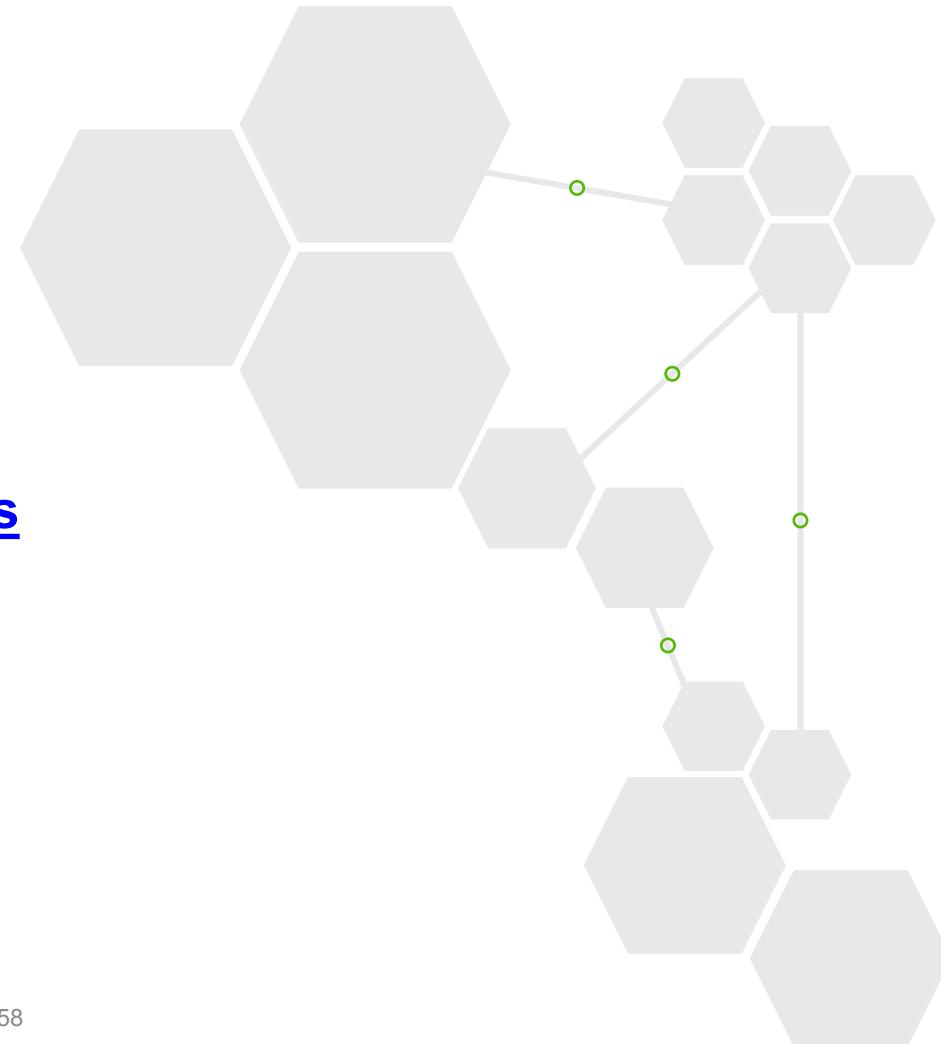
In-browser tutorials

my.storageos.com/main/tutorials

Quickstart

storageos.com/kubernetes

Hiring in London!





Questions

Slides at oicheryl.com

